

Notice that by storing A column by column in $A.band$, we obtain a column-oriented saxpy procedure. Indeed, Algorithm 1.2.2 is derived from Algorithm 1.1.4 by recognizing that each saxpy involves a vector with a small number of nonzeros. Integer arithmetic is used to identify the location of these nonzeros. As a result of this careful zero/nonzero analysis, the algorithm involves just $2n(p+q+1)$ flops with the assumption that p and q are much smaller than n .

1.2.6 Working with Diagonal Matrices

Matrices with upper and lower bandwidth zero are *diagonal*. If $D \in \mathbb{R}^{m \times n}$ is diagonal, then we use the notation

$$D = \text{diag}(d_1, \dots, d_q), \quad q = \min\{m, n\} \iff d_i = d_{ii}.$$

Shortcut notations when the dimension is clear include $\text{diag}(d)$ and $\text{diag}(d_i)$. Note that if $D = \text{diag}(d) \in \mathbb{R}^{n \times n}$ and $x \in \mathbb{R}^n$, then $Dx = d \cdot x$. If $A \in \mathbb{R}^{m \times n}$, then pre-multiplication by $D = \text{diag}(d_1, \dots, d_m) \in \mathbb{R}^{m \times m}$ scales rows,

$$B = DA \iff B(i, :) = d_i \cdot A(i, :), \quad i = 1:m$$

while post-multiplication by $D = \text{diag}(d_1, \dots, d_n) \in \mathbb{R}^{n \times n}$ scales columns,

$$B = AD \iff B(:, j) = d_j \cdot A(:, j), \quad j = 1:n.$$

Both of these special matrix-matrix multiplications require mn flops.

1.2.7 Symmetry

A matrix $A \in \mathbb{R}^{n \times n}$ is *symmetric* if $A^T = A$ and *skew-symmetric* if $A^T = -A$. Likewise, a matrix $A \in \mathbb{C}^{n \times n}$ is *Hermitian* if $A^H = A$ and *skew-Hermitian* if $A^H = -A$. Here are some examples:

$$\begin{array}{ll} \text{Symmetric:} & \begin{bmatrix} 1 & 2 & 3 \\ 2 & 4 & 5 \\ 3 & 5 & 6 \end{bmatrix}, \quad \text{Hermitian:} & \begin{bmatrix} 1 & 2-3i & 4-5i \\ 2+3i & 6 & 7-8i \\ 4+5i & 7+8i & 9 \end{bmatrix}, \\ \text{Skew-Symmetric:} & \begin{bmatrix} 0 & -2 & 3 \\ 2 & 0 & -5 \\ -3 & 5 & 0 \end{bmatrix}, \quad \text{Skew-Hermitian:} & \begin{bmatrix} i & -2+3i & -4+5i \\ 2+3i & 6i & -7+8i \\ 4+5i & 7+8i & 9i \end{bmatrix}. \end{array}$$

For such matrices, storage requirements can be halved by simply storing the lower triangle of elements, e.g.,

$$A = \begin{bmatrix} 1 & 2 & 3 \\ 2 & 4 & 5 \\ 3 & 5 & 6 \end{bmatrix} \iff A.vec = [1 \ 2 \ 3 \ 4 \ 5 \ 6].$$

For general n , we set

$$A.vec((n-j/2)(j-1)+i) = a_{ij} \quad 1 \leq j \leq i \leq n. \quad (1.2.2)$$

Here is a column-oriented gaxpy with the matrix A represented in $A.vec$.

Algorithm 1.2.3 (Symmetric Storage Gaxpy) Suppose $A \in \mathbb{R}^{n \times n}$ is symmetric and stored in the $A.vec$ style (1.2.2). If $x, y \in \mathbb{R}^n$, then this algorithm overwrites y with $y + Ax$.

```

for j = 1:n
    for i = 1:j - 1
        y(i) = y(i) + A.vec((i - 1)n - i(i - 1)/2 + j)x(j)
    end
    for i = j:n
        y(i) = y(i) + A.vec((j - 1)n - j(j - 1)/2 + i)x(j)
    end
end
end

```

This algorithm requires the same $2n^2$ flops that an ordinary gaxpy requires.

1.2.8 Permutation Matrices and the Identity

We denote the n -by- n identity matrix by I_n , e.g.,

$$I_4 = \begin{bmatrix} 1 & 0 & 0 & 0 \\ 0 & 1 & 0 & 0 \\ 0 & 0 & 1 & 0 \\ 0 & 0 & 0 & 1 \end{bmatrix}.$$

We use the notation e_i to designate the i th column of I_n . If the rows of I_n are reordered, then the resulting matrix is said to be a *permutation matrix*, e.g.,

$$P = \begin{bmatrix} 0 & 1 & 0 & 0 \\ 0 & 0 & 0 & 1 \\ 0 & 0 & 1 & 0 \\ 1 & 0 & 0 & 0 \end{bmatrix}. \quad (1.2.3)$$

The representation of an n -by- n permutation matrix requires just an n -vector of integers whose components specify where the 1's occur. For example, if $v \in \mathbb{R}^n$ has the property that v_i specifies the column where the "1" occurs in row i , then $y = Px$ implies that $y_i = x_{v_i}$, $i = 1:n$. In the example above, the underlying v -vector is $v = [2 \ 4 \ 3 \ 1]$.

1.2.9 Specifying Integer Vectors and Submatrices

For permutation matrix work and block matrix manipulation (§1.3) it is convenient to have a method for specifying structured integer vectors of subscripts. The MATLAB colon notation is again the proper vehicle and a few examples suffice to show how it works. If $n = 8$, then

$$\begin{aligned}
 v = 1:2:n & \implies v = [1 \ 3 \ 5 \ 7], \\
 v = n:-1:1 & \implies v = [8 \ 7 \ 6 \ 5 \ 4 \ 3 \ 2 \ 1], \\
 v = [(1:2:n) \ (2:2:n)] & \implies v = [1 \ 3 \ 5 \ 7 \ 2 \ 4 \ 6 \ 8].
 \end{aligned}$$